

Project: Mini-Project

Milan Lagae

Institution/University Name:

Faculty:

Course:

Vrije Universiteit Brussel

Sciences and Bioengineering Sciences

GPU Computing

Contents

1	Intro	1
2	Structure	1
3	Methodology	1
4	Part 1: Addition vs Multiplication	2
4.1	Setup	2
4.2	Memory Bandwidth	2
4.3	Compute Throughput	2
5	Part 2: Elements Per Thread	3
5.1	Setup	3
5.2	Memory Bandwidth	3
5.3	Time	3
6	Part 3: Roofline Model	4
7	Part 4: Local vs Global	4
8	Extra: Desktop vs Macbook	4
9	Appendix	5
9.1	Part 1	5
9.2	Part 2	6
9.3	Specifications	6
9.3.1	Macbook	6
9.3.2	Desktop	7
	Bibliography	8

1 Intro

Throughout all parts of the mini-project, it was chosen to compare two **operations** namely: addition & multiplication. To start, a comparison between the operations will be done, as requested in part 1. Following that, the kernel will be updated to execute on multiple elements simultaneous (one-to-many) mapping.

For the third part, the compute intensity will be increased by applying a loop count factor over the computations.

The final part will compare the local vs global memory location, as indicated in the scheme¹ in the assignment.

2 Structure

TODO

3 Methodology

For each iteration of the combination of values to be measures, **15** runs where measured. The first 5 runs where discarded to allow the system to stabilize. This is in accordance to the recommendations in [1]. For each measured value, is applicable a Cov range chart is produced and will be reference, these can be found in §9. The acceptable range of the Cov is taken from [2].

¹http://parallel.vub.ac.be/education/gpu/labs/miniproject_step4_with_local%20memory.jpg

4 Part 1: Addition vs Multiplication

4.1 Setup

TODO

4.2 Memory Bandwidth




Figure 1:

TODO




Figure 2:

4.3 Compute Throughput




Figure 3:

5 Part 2: Elements Per Thread

5.1 Setup

The array size is fixed for all benchmark variations performed in this part to 2^{20} .

5.2 Memory Bandwidth




Figure 4:




Figure 5:

5.3 Time




Figure 6:




Figure 7:

- 6 Part 3: Roofline Model**
- 7 Part 4: Local vs Global**
- 8 Extra: Desktop vs Macbook**

9 Appendix

9.1 Part 1




Figure 8:




Figure 9:

9.2 Part 2




Figure 10:




Figure 11:

9.3 Specifications

9.3.1 Macbook

Part	Value
CPU	M2 Pro (6 performance and 4 efficiency)
RAM	16GB
OS	TODO

Table 1: Macbook Specifications

9.3.2 Desktop

Part	Value
CPU	Ryzen 9 5950X
GPU	RTX 3070
OpenCL	TODO
Driver Version	TODO
RAM	64GB (3200 Mhz)
OS	Windows Version 10.0.22631 Build 22631

Table 2: Desktop Specifications

Bibliography

- [1] P. A. Paolillo, “Lecture 2 - Metrics, Distributions & Experimental Rigor - How Many runs do I need?.” p. 38, 2025.
- [2] P. A. Paolillo, “Lecture 2 - Metrics, Distributions & Experimental Rigor - Coefficient of Variation (CoV).” p. 25, 2025.